

# in3050\_2023\_w7\_solutions

February 26, 2023

## 1 Week 07: Linear and logistic regression

### 1.0.1 Introduction

This week, we will get some first-hand experience with regression. We will implement gradient descent for linear regression. Then we will proceed to classification, first by using linear regression and then logistic regression.

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import sklearn
```

### 1.1 NumPy

We will first familiarize ourselves a little with NumPy. A function which we will use over again is `linspace(x1,x2,N)` which makes a vector of length  $N$  splitting the interval  $[x1, x2]$  into equally sized intervals.

```
[2]: xx = np.linspace(-5,5,101)
xx[:10]
```

```
[2]: array([-5. , -4.9, -4.8, -4.7, -4.6, -4.5, -4.4, -4.3, -4.2, -4.1])
```


One of the major improvements from using NumPy is the possibility of computing many values by applying a function to a numpy array.

```
[3]: y1 = -6*xx**3 + xx**2 -3*xx + 5
```

`y1` will contain the corresponding function values for each element `x` in `xx`. We may plot the result.

```
[4]: plt.plot(xx,y1)
```

```
[4]: [<matplotlib.lines.Line2D at 0x1975b525240>]
```




## 1.2 Dataset for linear regression

We will start with a smooth curve and add some “noise”. The underlying idea is that the smooth curve represents the function we are looking for, and that this is the best we can hope to learn. A solution which does better on the training material than the smooth curve is probably overfit and will not generalize as well to new data as the smooth curve. We are using a normal distribution to generate noise. The numpy function `normal` will generate a vector of `size` many random points around `loc` from a distribution with standard deviation `scale`.

```
[5]: from numpy.random import normal
     t = y1 + normal(loc=0, scale=50, size=101)
     plt.scatter(xx, t)
```

```
[5]: <matplotlib.collections.PathCollection at 0x1975be1eda0>
```



Our training data now consists of pairs  $(xx[i], t[i])$ , where  $xx[i]$  is the datapoint and  $t[i]$  the target value. So far, both  $xx$  and  $t$  are vectors. Check their shapes, e.g., `xx.shape`. The goal is to make an implementation for linear regression which works with an arbitrary number of input features and not just one. We will therefore transform  $xx$  to a matrix of dimension  $N \times m$  where each row represents one datapoint, and  $m$  is the number of input variables (or features). Check the shape of  $X$  after the transform.

```
[6]: X = xx.reshape(-1,1)
```

#### **Solution**

```
[7]: xx.shape
```

```
[7]: (101,)
```

```
[8]: X.shape
```

```
[8]: (101, 1)
```

**End of solution**

### 1.3 Part A: Linear regression

We will implement our own linear regression model. Our aim is to find an approximate function that fits the data generated above.

Since we are dealing with only one input variable, we start with a simple linear function,  $f(x_1) = w_0 + w_1x_1$ .

#### 1.3.1 Exercise 1: MSE

We wonder if our  $f$  fits the data well, and what parameters will give us the best approximation. We will estimate this using the Mean Squared Error:

$$\frac{1}{N} \sum_{j=1}^N (t_j - \sum_{i=0}^m w_i x_{ji})^2$$

Write a function calculating MSE of our approximation.

```
[9]: # Your code here
```

**solution**

```
[10]: # Pure python solution
def mse_1(y, y_pred):
    sum_errors = 0.
    for i in range(0, len(y)):
        sum_errors += (y[i] - y_pred[i])**2
    mean_squared_error = sum_errors/len(y)
    return mean_squared_error
```

```
[11]: # Vector form solution
def mse_2(x, y):
    return np.mean((x - y)**2)
```

**end of solution**

#### 1.3.2 Exercise 2: Testing the MSE

To test our implementation, we can take the function  $f(x_1) = 0$  as a baseline and calculate the MSE for this  $f$ . Also calculate the Root Means Square Error which provides a more natural measure for how good the fit is.

**Solution**

```
[12]: hypothesis = np.array([0 for x in xx])
mse = mse_1(hypothesis, t)
print("MSE: ", mse)
```

MSE: 92493.0749194188

```
[13]: hypothesis = np.array([0 for x in xx])
mse = mse_2(hypothesis, t)
```

```
print("MSE: ", mse)
```

MSE: 92493.0749194188

```
[14]: print("RMSE: ", np.sqrt(mse))
```

RMSE: 304.1267415394753

End of solution

### 1.3.3 Exercise 3: Adding bias

We will implement linear regression with gradient descent and test it on the data. To make it simple, we will add a  $x_0 = 1$  to all our datapoints, and consider  $f(x_1) = w_0 + w_1x_1$  as  $f(x_0, x_1) = w_0x_0 + w_1x_1$ . Make a procedure that does this.

```
[15]: def add_bias(X):  
    """X is a Nxm matrix: N datapoints, m features  
    Return a Nx(m+1) matrix with added bias in position zero"""  
    pass
```

**solution** There are many ways that can be done in numpy. This is one of them.

```
[16]: def add_bias(X):  
    """X is a Nxm matrix: N datapoints, m features  
    Return a Nx(m+1) matrix with added bias in position zero"""  
    N = X.shape[0]  
    bias = np.ones((N,1)) # Make a N*1 matrix of 1-s  
    # Concatenate the column of bias in front of the columns of X.  
    return np.concatenate((bias, X), axis = 1)
```

end of solution

### 1.3.4 Exercise 4: Gradient Descent

We will implement the linear regression in a class as we did with the classifiers earlier. The fit method will run the gradient descent step a number of times to train the classifier. The predict method should take a matrix containing several data points and predict the outcome for all of them. Fill in the methods.

Assume that the matrix of training data is not extended with bias features. Hence, make adding bias a part of your methods.

After training there should be an attribute with learned coefficients (weights) which is applied by the predict method.

```
[17]: class NumpyLinReg():  
  
    def fit(self, X_train, t_train, eta = 0.1, epochs=10):
```

```

"""X_train is a Nxm matrix, N data points, m features
t_train is a vector of length N,
the targets values for the training data"""

```

```

def predict(self, X):
    """X is a Kxm matrix for some K>=1
    predict the value for each point in X"""
    pass

```

solution

```

[18]: class NumpyLinReg():

    def fit(self, X_train, t_train, eta = 0.1, epochs=10):
        """X_train is a Nxm matrix, N data points, m features
        t_train is a vector of length N,
        the targets values for the training data"""

        (N, m) = X_train.shape
        X_train = add_bias(X_train)

        self.weights = weights = np.zeros(m+1)

        for e in range(epochs):
            weights -= eta / N * X_train.T @ (X_train @ weights - t_train)

    def predict(self, X):
        """X is a Kxm matrix for some K>=1
        predict the value for each point in X"""
        Z = add_bias(X)
        return Z @ self.weights

```

end of solution

### 1.3.5 Exercise 5: Train and test the model

Fit the model to the training data. Report the coefficients. Plot the line together with the observations. Calculate the RMSE. Is the result a better fit than the baseline constant function  $f(x) = 0$ ?

solution

```

[19]: reg = NumpyLinReg()
      reg.fit(X,t, epochs=100)
      print("The coefficients: ", reg.weights)


```

The coefficients: [ 11.43131192 -95.77449346]

```
[20]: plt.plot(X, reg.weights[0]+reg.weights[1]*X, color="r")
plt.plot(X, [0 for i in X], 'g')

# The generated dataset
plt.scatter(X, t)
```

[20]: <matplotlib.collections.PathCollection at 0x1975becb640>



```
[21]: error = mse_2(t, add_bias(X) @ reg.weights)
print("MSE: ", error)
```

MSE: 14393.987508996268

```
[22]: print("RMSE :", np.sqrt(error))
```

RMSE : 119.9749453385842

It has improved

end of solution


## 1.4 Dataset for classification

We will use simple synthetic data similarly to week\_05, but we will make the set a little bigger to get more reliable results.

```
[23]: from sklearn.datasets import make_blobs
X_train, y_train = make_blobs(n_samples=500, centers=[[0,0],[1,2]],
                             n_features=2, random_state=2019)
X_test, y_test = make_blobs(n_samples=500, centers=[[0,0],[1,2]],
                             n_features=2, random_state=2020)
```

```
[24]: def show(X, y, marker='.'):
    labels = set(y)
    for lab in labels:
        plt.plot(X[y == lab][:, 1], X[y == lab][:, 0],
                 marker, label="class {}".format(lab))
    plt.legend()
```

```
[25]: show(X_train, y_train)
```





## 1.5 Linear Regression classifier

This is also called Ridge Classifier in the literature when it is smoothed. We will consider the simple unsmoothed version here and return to smoothing and regularization in a later lecture.

### 1.5.1 Exercise 6: Coding the classifier

Make a linear regression classifier. Make it as a Python class with methods for ‘fit’ and ‘predict’, similarly to the linear regression above and the  $k$ NN classifier from week05 exercises.

#### Solution

```
[27]: class NumpyClassifier():
        """Common methods to all numpy classifiers --- if any"""

[28]: class NumpyLinRegClass(NumpyClassifier):

        def fit(self, X_train, t_train, eta = 0.1, epochs=10):
            """X_train is a Nxm matrix, N data points, m features
            t_train is a vector of length N,
            the targets values for the training data"""

            (N, m) = X_train.shape
            X_train = add_bias(X_train)

            self.weights = weights = np.zeros(m+1)

            for e in range(epochs):
                weights -= eta / N * X_train.T @ (X_train @ weights - t_train)

        def predict(self, x, threshold=0.5):
            """X is a Kxm matrix for some K>=1
            predict the value for each point in X"""
            z = add_bias(x)
            score = z @ self.weights
            return score>threshold
```

End of solution

### 1.5.2 Exercise 7: Experiment

We will conduct repeated testing. We therefore need a development test set different from the final test set. Make such a set  $X_{dev}$ ,  $y_{dev}$ , similarly to  $X_{test}$ ,  $y_{test}$  using `random_state=2021`. Train the classifier on  $X_{train}$ ,  $y_{train}$  and test for accuracy on  $X_{dev}$ ,  $y_{dev}$ .

You also need a procedure for accuracy here.

#### Solution

```
[29]: X_dev, y_dev = make_blobs(n_samples=500, centers=[[0,0],[1,2]],
                                n_features=2, random_state=2021)
```

```
[30]: # The week05 solution

def accuracy_1(y, t):
    """Calculate the accuracy"""
    equal = len([(p, g) for (p,g) in zip(y, t) if p==g])
    return equal / len(t)
```

```
[31]: # Vector solution

def accuracy(y, t):
    return np.mean(y == t)
```

```
[32]: lin_cl = NumpyLinRegClass()
lin_cl.fit(X_train, y_train)
accuracy(lin_cl.predict(X_dev), y_dev)
```

```
[32]: 0.864
```

End of solution

## 1.6 Logistic Regression

### 1.6.1 Exercise 8: The logistic function

Implement the logistic function. Sometimes called only the sigmoid.

```
[33]: def logistic(x):
        # fill in the rest
        pass
```

Solution

```
[34]: def logistic(x):
        return 1/(1+np.exp(-x))
```

End of solution

### 1.6.2 Exercise 9: Code for the classifier

Write code for the logistic regression classifier. Compared to linear regression classifier you have to make adaptations to both fit and predict taking the logistic into consideration

Solution

```
[35]: class NumpyLogReg(NumpyClassifier):
```

```

def fit(self, X_train, t_train, eta = 0.1, epochs=10):
    """X_train is a Nxm matrix, N data points, m features
    t_train is a vector of length N,
    the targets values for the training data"""

    (N, m) = X_train.shape
    X_train = add_bias(X_train)

    self.weights = weights = np.zeros(m+1)

    for e in range(epochs):
        weights -= eta / N * X_train.T @ (self.forward(X_train) - t_train)

def forward(self, X):
    return logistic(X @ self.weights)

def predict(self, x, threshold=0.5):
    """X is a Kxm matrix for some K>=1
    predict the value for each point in X"""
    z = add_bias(x)
    return (self.forward(z) > threshold).astype('int')

```

End of solution

### 1.6.3 Exercise 10: Initial experiments

Train the classifier on  $X_{\text{train}}$ ,  $y_{\text{train}}$  and test for accuracy on  $X_{\text{dev}}$ ,  $y_{\text{dev}}$ .

**Solution**

```

[36]: lr_cl = NumpyLogReg()
      lr_cl.fit(X_train, y_train)
      accuracy(lr_cl.predict(X_dev), y_dev)

```

[36]: 0.752

End of solution

### 1.6.4 Exercise 11: Repeated experimentation

Did you get better results than with the linear regression classifier? That does not necessarily have to be the case for this data set. But, if your result is much inferior to the linear regression classifier, the reason might be the parameter settings. Experiment with the parameter values for the learning rate and the number of epochs to get an optimal result.

**Solution**

```
[37]: for e in [1, 2, 5, 10, 50, 100, 1000, 10000, 100000]:
    lr_cl = NumpyLogReg()
    lr_cl.fit(X_train, y_train, epochs=e, eta=0.1)
    print("Learning rate: {} Epochs: {:7} Accuracy: {}".format(
        0.1, e, accuracy(lr_cl.predict(X_dev), y_dev)))
```

```
Learning rate: 0.1 Epochs:      1 Accuracy: 0.74
Learning rate: 0.1 Epochs:      2 Accuracy: 0.742
Learning rate: 0.1 Epochs:      5 Accuracy: 0.744
Learning rate: 0.1 Epochs:     10 Accuracy: 0.752
Learning rate: 0.1 Epochs:     50 Accuracy: 0.818
Learning rate: 0.1 Epochs:    100 Accuracy: 0.842
Learning rate: 0.1 Epochs:   1000 Accuracy: 0.87
Learning rate: 0.1 Epochs:  10000 Accuracy: 0.874
Learning rate: 0.1 Epochs: 100000 Accuracy: 0.874
```

```
[38]: for eta in [1, 0.1, 0.01, 0.001, 0.0001, 0.00001]:
    lr_cl = NumpyLogReg()
    lr_cl.fit(X_train, y_train, epochs=1000, eta = eta)
    print("Learning rate: {:7} Epochs: {:4} Accuracy: {}".format(
        eta, 1000, accuracy(lr_cl.predict(X_dev), y_dev)))
```

```
Learning rate:      1 Epochs: 1000 Accuracy: 0.874
Learning rate:    0.1 Epochs: 1000 Accuracy: 0.87
Learning rate:   0.01 Epochs: 1000 Accuracy: 0.842
Learning rate:  0.001 Epochs: 1000 Accuracy: 0.76
Learning rate: 0.0001 Epochs: 1000 Accuracy: 0.742
Learning rate: 1e-05 Epochs: 1000 Accuracy: 0.74
```

We see that the best accuracy is 0.874. We also see that it is faster to reach the solution with a learning rate of 1.0 than with smaller learning rates. ##### End of solution

## 1.7 End of week 07